

ml ex 3 submit

lrCostFunction

```
function [J, grad] = lrCostFunction(theta, X, y, lambda)
%LRCOSTFUNCTION Compute cost and gradient for logistic regression with
%regularization
% J = LRCOSTFUNCTION(theta, X, y, lambda) computes the cost of using
% theta as the parameter for regularized logistic regression and the
% gradient of the cost w.r.t. to the parameters.

% Initialize some useful values
m = length(y); % number of training examples

% You need to return the following variables correctly
J = 0;
grad = zeros(size(theta));

% ===== YOUR CODE HERE =====
% Instructions: Compute the cost of a particular choice of theta.
%             You should set J to the cost.
%             Compute the partial derivatives and set grad to the partial
%             derivatives of the cost w.r.t. each parameter in theta
%
% Hint: The computation of the cost function and gradients can be
%       efficiently vectorized. For example, consider the computation
%
%       sigmoid(X * theta)
%
%       Each row of the resulting matrix will contain the value of the
%       prediction for that example. You can make use of this to vectorize
%       the cost function and gradient computations.
%
% Hint: When computing the gradient of the regularized cost function,
%       there're many possible vectorized solutions, but one solution
%       looks like:
%       grad = (unregularized gradient for logistic regression)
%       temp = theta;
%       temp(1) = 0; % because we don't add anything for j = 0
%       grad = grad + YOUR_CODE_HERE (using the temp variable)
%
```

```

h = sigmoid(X*theta);
reg = (lambda / (2 * m)) * (sum(theta.^2) - theta(1)^2);
J = (-1 / m) * sum((y .* log(h)) + (1 - y) .* log(1 - h)) + reg;

```

```

temp = (h - y)' * X;
grad = (1 / m) * temp' + (lambda / m) .* theta;
grad(1) = (1 / m) * ((h - y)' * X(:, 1));

```

```

%

```

```

H = sigmoid(X*theta);
T = y.*log(H) + (1 - y).*log(1 - H);
J = -1/m*sum(T) + lambda/(2*m)*sum(theta(2:end).^2);

```

```

ta = [0; theta(2:end)];
grad = X'*(H - y)/m + lambda/m*ta;

```

```

% =====

```

```

=====

```

```

grad = grad(:);

```

```

end

```

```

OneVsAll

```

```

function [all_theta] = oneVsAll(X, y, num_labels, lambda)
%ONEVSALL trains multiple logistic regression classifiers and returns all
%the classifiers in a matrix all_theta, where the i-th row of all_theta
%corresponds to the classifier for label i
% [all_theta] = ONEVSALL(X, y, num_labels, lambda) trains num_labels
% logistic regression classifiers and returns each of these classifiers
% in a matrix all_theta, where the i-th row of all_theta corresponds
% to the classifier for label i

```

```

% Some useful variables
m = size(X, 1);
n = size(X, 2);

```

```

% You need to return the following variables correctly
all_theta = zeros(num_labels, n + 1);

```

```

% Add ones to the X data matrix

```

```
X = [ones(m, 1) X];
```

```
% ===== YOUR CODE HERE =====
```

```
% Instructions: You should complete the following code to train num_labels
```

```
%     logistic regression classifiers with regularization
```

```
%     parameter lambda.
```

```
%
```

```
% Hint: theta(:) will return a column vector.
```

```
%
```

```
% Hint: You can use y == c to obtain a vector of 1's and 0's that tell you
```

```
%     whether the ground truth is true/false for this class.
```

```
%
```

```
% Note: For this assignment, we recommend using fmincg to optimize the cost
```

```
%     function. It is okay to use a for-loop (for c = 1:num_labels) to
```

```
%     loop over the different classes.
```

```
%
```

```
%     fmincg works similarly to fminunc, but is more efficient when we
```

```
%     are dealing with large number of parameters.
```

```
%
```

```
% Example Code for fmincg:
```

```
%
```

```
%     % Set Initial theta
```

```
%     initial_theta = zeros(n + 1, 1);
```

```
%
```

```
%     % Set options for fminunc
```

```
%     options = optimset('GradObj', 'on', 'MaxIter', 50);
```

```
%
```

```
%     % Run fmincg to obtain the optimal theta
```

```
%     % This function will return theta and the cost
```

```
%     [theta] = ...
```

```
%         fmincg (@(t)(lrCostFunction(t, X, (y == c), lambda)), ...
```

```
%             initial_theta, options);
```

```
%
```

```
for c = 1:num_labels
```

```
    initial_theta = zeros(n + 1, 1);
```

```
    options = optimset('GradObj', 'on', 'maxIter', 50);
```

```
    [initial_theta] = fmincg (@(t)(lrCostFunction(t, X, (y == c), lambda)),  
initial_theta, options);
```

```
    all_theta(c,:) = initial_theta;
```

```
% =====
```

```
for c = 1 : num_labels,
```

```
    initial_theta = zeros(n + 1 , 1);
```

```
    options = optimset('GradObj' , 'on' , 'MaxIter' , 50);
```

```

        [theta] = ...
            fmincg(@(t)(lrCostFunction(t , X , (y == c) , lambda)), ...
                initial_theta , options);
        all_theta(c,:) = theta';
    end
end

```

predictOneVsAll

```

function p = predictOneVsAll(all_theta, X)
%PREDICT Predict the label for a trained one-vs-all classifier. The labels
%are in the range 1..K, where K = size(all_theta, 1).
% p = PREDICTONEVSALL(all_theta, X) will return a vector of predictions
% for each example in the matrix X. Note that X contains the examples in
% rows. all_theta is a matrix where the i-th row is a trained logistic
% regression theta vector for the i-th class. You should set p to a vector
% of values from 1..K (e.g., p = [1; 3; 1; 2] predicts classes 1, 3, 1, 2
% for 4 examples)

```

```

m = size(X, 1);
num_labels = size(all_theta, 1);

```

```

% You need to return the following variables correctly
p = zeros(size(X, 1), 1);

```

```

% Add ones to the X data matrix
X = [ones(m, 1) X];

```

```

% ===== YOUR CODE HERE =====
% Instructions: Complete the following code to make predictions using
%             your learned logistic regression parameters (one-vs-all).
%             You should set p to a vector of predictions (from 1 to
%             num_labels).
%
% Hint: This code can be done all vectorized using the max function.
%       In particular, the max function can also return the index of the
%       max element, for more information see 'help max'. If your examples
%       are in rows, then, you can use max(A, [], 2) to obtain the max
%       for each row.
%

```

```

for c = 1:m
    one_example = X(c,:);
    [max_value, max_index] = max(all_theta * one_example');
    p(c,1) = max_index;

```

```

%

```

```

for c = 1 : num_labels,
    initial_theta = zeros(n + 1 , 1);
    options = optimset('GradObj' , 'on' , 'MaxIter' , 50);
    [theta] = ...
        fmincg(@(t)(lrCostFunction(t , X , (y == c) , lambda)), ...
            initial_theta , options);
    all_theta(c,:) = theta';

=====
end

predict

function p = predict(Theta1, Theta2, X)
%PREDICT Predict the label of an input given a trained neural network
% p = PREDICT(Theta1, Theta2, X) outputs the predicted label of X given the
% trained weights of a neural network (Theta1, Theta2)

% Useful values
m = size(X, 1);
num_labels = size(Theta2, 1);

% You need to return the following variables correctly
p = zeros(size(X, 1), 1);

% ===== YOUR CODE HERE =====
% Instructions: Complete the following code to make predictions using
% your learned neural network. You should set p to a
% vector containing labels between 1 to num_labels.
%
% Hint: The max function might come in useful. In particular, the max
% function can also return the index of the max element, for more
% information see 'help max'. If your examples are in rows, then, you
% can use max(A, [], 2) to obtain the max for each row.
%

X = [ones(m, 1) X];

for c = 1:m
    one_example = X(c,:);
    z2 = Theta1 * one_example';
    a2 = sigmoid(z2);

    a2 = [ones(1, 1); a2];

    z3 = Theta2 * a2;

```

```
a3 = sigmoid(z3);

[max_value, max_index] = max(a3);
p(c, 1) = max_index;

%
C = sigmoid(X*all_theta');
[M , p] = max(C , [] , 2);

=====
end
```