

ml ex5 submit

```
function [J, grad] = linearRegCostFunction(X, y, theta, lambda)
    %LINEARREGCOSTFUNCTION Compute cost and gradient for regularized
linear
    %regression with multiple variables
    % [J, grad] = LINEARREGCOSTFUNCTION(X, y, theta, lambda) computes
the
    % cost of using theta as the parameter for linear regression to fit the
    % data points in X and y. Returns the cost in J and the gradient in grad

    % Initialize some useful values
    m = length(y); % number of training examples

    % You need to return the following variables correctly
    J = 0;
    grad = zeros(size(theta));

    % ===== YOUR CODE HERE =====
    % Instructions: Compute the cost and gradient of regularized linear
    % regression for a particular choice of theta.
    %
    % You should set J to the cost and grad to the gradient.
    %

    H = X*theta;
    J = sum((H - y).^2) / (2 * m) + lambda*sum(theta(2:end).^2) / (2 * m);

    grad = X'*(H - y) / m + lambda*[0;theta(2:end)] / m;

    % =====
    grad = grad(:);

end

function [error_train, error_val] = ...
    learningCurve(X, y, Xval, yval, lambda)
    %LEARNINGCURVE Generates the train and cross validation set errors
needed
    %to plot a learning curve
    % [error_train, error_val] = ...
```

```

% LEARNINGCURVE(X, y, Xval, yval, lambda) returns the train and
% cross validation set errors for a learning curve. In particular,
% it returns two vectors of the same length - error_train and
% error_val. Then, error_train(i) contains the training error for
% i examples (and similarly for error_val(i)).
%
% In this function, you will compute the train and test errors for
% dataset sizes from 1 up to m. In practice, when working with larger
% datasets, you might want to do this in larger intervals.
%
% Number of training examples
m = size(X, 1);

% You need to return these values correctly
error_train = zeros(m, 1);
error_val = zeros(m, 1);

% ===== YOUR CODE HERE =====
% Instructions: Fill in this function to return training errors in
% error_train and the cross validation errors in error_val.
% The vector numex_vec contains the number of training
% examples to use for each calculation of training error and
% cross validation error, i.e, error_train(i) and
% error_val(i) should give you the errors
% obtained after training on i examples.
%
% Note: You should evaluate the training error on the first i training
% examples (i.e., X(1:i, :) and y(1:i)).
%
% For the cross-validation error, you should instead evaluate on
% the _entire_ cross validation set (Xval and yval).
%
% Note: If you are using your cost function (linearRegCostFunction)
% to compute the training and cross validation error, you should
% call the function with the lambda argument set to 0.
% Do note that you will still need to use lambda when running
% the training to obtain the theta parameters.
%
% Hint: You can loop over the examples with the following:
%
% for i = 1:m
%     % Compute train/cross validation errors using training examples
%     % X(1:i, :) and y(1:i), storing the result in
%     % error_train(i) and error_val(i)
%     ....

```

```

%
%     end
%
% ----- Sample Solution -----
val_size = size(Xval,1);

for i = 1 : m,
    [theta] = trainLinearReg([ones(i , 1) X(1:i , :)], y(1:i), lambda);
    [error_train(i), grad] = linearRegCostFunction([ones(i , 1) X(1:i , :)],
y(1:i), theta, 0);
    [error_val(i), grad] = linearRegCostFunction([ones(val_size , 1) Xval],
yval, theta, 0);
end
% -----
% =====
end

function [X_poly] = polyFeatures(X, p)
    %POLYFEATURES Maps X (1D vector) into the p-th power
    % [X_poly] = POLYFEATURES(X, p) takes a data matrix X (size m x 1) and
    % maps each example into its polynomial features where
    % X_poly(i, :) = [X(i) X(i).^2 X(i).^3 ... X(i).^p];
    %
    % You need to return the following variables correctly.
    X_poly = zeros(numel(X), p);

    % ===== YOUR CODE HERE =====
    % Instructions: Given a vector X, return a matrix X_poly where the p-th
    % column of X contains the values of X to the p-th power.
    %
    %
    %

    for i = 1 : p,
        X_poly(:,i) = X.^i;

    % =====

function [lambda_vec, error_train, error_val] = ...
    validationCurve(X, y, Xval, yval)
    %VALIDATIONCURVE Generate the train and validation errors needed to
    %plot a validation curve that we can use to select lambda
    % [lambda_vec, error_train, error_val] = ...
    % VALIDATIONCURVE(X, y, Xval, yval) returns the train
    % and validation errors (in error_train, error_val)
    % for different values of lambda. You are given the training set (X,
    % y) and validation set (Xval, yval).

```

```

%
% Selected values of lambda (you should not change this)
lambda_vec = [0 0.001 0.003 0.01 0.03 0.1 0.3 1 3 10]';
% You need to return these variables correctly.
error_train = zeros(length(lambda_vec), 1);
error_val = zeros(length(lambda_vec), 1);

% ===== YOUR CODE HERE =====
% Instructions: Fill in this function to return training errors in
%           error_train and the validation errors in error_val. The
%           vector lambda_vec contains the different lambda parameters
%           to use for each calculation of the errors, i.e,
%           error_train(i), and error_val(i) should give
%           you the errors obtained after training with
%           lambda = lambda_vec(i)
%
% Note: You can loop over lambda_vec with the following:
%
%   for i = 1:length(lambda_vec)
%       lambda = lambda_vec(i);
%       % Compute train / val errors when training linear
%       % regression with regularization parameter lambda
%       % You should store the result in error_train(i)
%       % and error_val(i)
%       ....
%   end
%
%
for i = 1 : length(lambda_vec),
    lambda = lambda_vec(i);
    [theta] = trainLinearReg(X, y, lambda);
    [error_train(i), grad] = linearRegCostFunction(X, y, theta, 0);
    [error_val(i), grad] = linearRegCostFunction(Xval , yval, theta, 0);
end
% =====
end

```