Week 2 Quiz

1.Question 1

What is a windowed dataset?

A fixed-size subset of a time series

2.Question 2

What does 'drop_remainder=true' do?

It ensures that all rows in the data window are the same length by cropping data

3.Question 3

What's the correct line of code to split an n column window into n-1 columns for features and 1 column for a label

dataset = dataset.map(lambda window: (window[:-1], window[-1:]))

4.Question 4

What does MSE stand for?

Mean Squared error

5.Question 5

What does MAE stand for?

Mean Absolute Error

6.Question 6

If time values are in time[], series values are in series[] and we want to split the

series into training and validation at time 1000, what is the correct code

time_train = time[:split_time]

x_train = series[:split_time]

time_valid = time[split_time:]

x_valid = series[split_time:]

7.Question 7

If you want to inspect the learned parameters in a layer after training, what's a good technique to use?

Assign a variable to the layer and add it to the model using that variable. Inspect its properties after training

8.Question 8

How do you set the learning rate of the SGD optimizer?

Use the lr property

9.Question 9

If you want to amend the learning rate of the optimizer on the fly, after each epoch, what do you do?

Use a LearningRateScheduler object in the callbacks namespace and assign that to the callback

END